

▼ Stock Market Prediction Using ARIMA in Python

▼ Finance Analytics Project

```
1 import pandas as pd
2 import numpy as np
3 import math
4
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7 from pylab import rcParams
8
9 import statsmodels.api as sm
10 from statsmodels.tsa.stattools import adfuller
11 from statsmodels.tsa.seasonal import seasonal_decompose
12 from statsmodels.tsa.arima.model import ARIMA
13 from pmdarima.arima import auto_arima
14
15 from sklearn.metrics import mean_squared_error, mean_absolute_error
```

```
1 df = pd.read_csv('HCLTECH.NS.csv')
2 df.head()
```

	Date	Open	High	Low	Close	Adj Close	Volume
0	2017-04-21	408.500000	410.725006	403.899994	406.375000	358.572723	1780310
1	2017-04-24	406.500000	412.000000	404.625000	409.774994	361.572754	1373968
2	2017-04-25	412.000000	412.924988	406.799988	409.975006	361.749268	2194602
3	2017-04-26	410.950012	411.000000	398.274994	400.125000	353.057892	1918248
4	2017-04-27	398.500000	410.000000	398.100006	404.850006	357.227112	5640334

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1235 entries, 0 to 1234
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
0   Date        1235 non-null   object
1   Open        1235 non-null   float64
2   High        1235 non-null   float64
3   Low         1235 non-null   float64
4   Close       1235 non-null   float64
5   Adj Close   1235 non-null   float64
6   Volume      1235 non-null   int64
dtypes: float64(5), int64(1), object(1)
memory usage: 67.7+ KB
```

▼ Data Cleaning and Transformation

▼ Removing duplicated date values if any

```
1 df[df["Date"].duplicated(keep = False)]
2 df = df[~df["Date"].duplicated()]
```

▼ Null check to identify missing values

```
1 df.isnull().sum()/len(df)
```

```
Date      0.0
Open      0.0
High      0.0
Low       0.0
Close     0.0
Adj Close 0.0
Volume    0.0
dtype: float64
```

Transforming 'Date' column to index

```
1 df.index = pd.to_datetime(df["Date"])
2 data = df[df.columns[1:]]
3 data.head()
```

	Open	High	Low	Close	Adj Close	Volume
Date						
2017-04-21	408.500000	410.725006	403.899994	406.375000	358.572723	1780310
2017-04-24	406.500000	412.000000	404.625000	409.774994	361.572754	1373968
2017-04-25	412.000000	412.924988	406.799988	409.975006	361.749268	2194602
2017-04-26	410.950012	411.000000	398.274994	400.125000	353.057892	1918248
2017-04-27	398.500000	410.000000	398.100006	404.850006	357.227112	5640334

Data Visualization

We consider 'Adj Close' column to predict the future stock price, and it looks like below:

```
1 plt.figure(figsize = (20, 10))
2 sns.set_style('darkgrid')
3 plt.xlabel('Date', fontsize = 20)
4 plt.ylabel('Close Price', fontsize = 20)
5 plt.title('HCL Stock Market Closing Price', fontsize = 20)
6 plt.plot(data['Close'])
7 plt.tick_params(axis = 'x', labels = 18)
8 plt.tick_params(axis = 'y', labels = 18)
9 plt.savefig('AdjClose_TimeSeriesPlot.png')
```



ADF Test to check if data is Stationary or not


```
1 def test_adf(timeseries):
2     moving_average = timeseries.rolling(12).mean()
3     moving_std = timeseries.rolling(12).std()
4     plt.figure(figsize = (20,10))
5     plt.plot(timeseries, color = 'blue', label = 'Original')
6     plt.plot(moving_average, color = 'red', label = 'Rolling Mean')
7     plt.plot(moving_std, color = 'black', label = 'Rolling Std')
8     plt.legend(loc = 'best', fontsize = 18)
```

```

9 plt.title('Rolling Mean and Standard Deviation', fontsize = 20)
10 plt.tick_params(axis = 'x', labelsiz = 18)
11 plt.tick_params(axis = 'y', labelsiz = 18)
12 plt.savefig('ADFTest2.png')
13 plt.show(block = False)
14 print("Results of Dicky Fuller Test")
15 adft = adfuller(timseries, autolag = 'AIC')
16 output = pd.Series(adft[0:4], index = ['Test Statistics', 'p-value', 'No. of lags used', 'Number of observations used'])
17 for key, value in adft[4].items():
18     output['Critical value (%s)' %key] = value
19 print(output)

```

```
1 test_adf(data['Close'])
```



```

Results of Dicky Fuller Test
Test Statistics          -0.684578
p-value                 0.850704
No. of lags used        7.000000
Number of observations used 1227.000000
Critical value (1%)     -3.435691
Critical value (5%)     -2.863898
Critical value (10%)    -2.568026
dtype: float64

```

From the above result, it is evident that test Statistics is greater than the critical values and $p > 0.05$, hence we fail to reject the null hypothesis meaning our data is non-stationary.

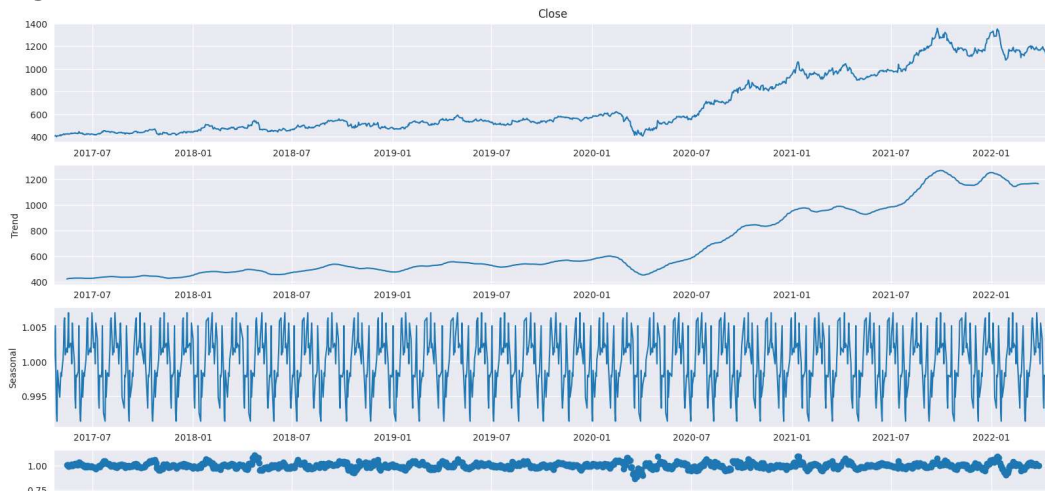
▼ From the time series data, let's separate trend and seasonality

```

1 result = seasonal_decompose(data['Close'], model = 'multiplicative', period = 30)
2 fig = plt.figure()
3 fig = result.plot()
4 fig.savefig('stationaryData.png')
5 fig.set_size_inches(16, 9)

```

<Figure size 2000x1000 with 0 Axes>




▼ Converting Non-stationary to Stationary Data

```

1 rcParams['figure.figsize'] = 20, 10
2 data_adj_close_log = np.log(data['Close'])
3 moving_average = data_adj_close_log.rolling(12).mean()
4 std_dev = data_adj_close_log.rolling(12).std()
5 plt.legend(loc = 'best')
6 plt.title('Moving Average', fontsize = 20)
7 plt.plot(std_dev, color = "black", label = "Standard Deviation")
8 plt.plot(moving_average, color = "red", label = "Mean")
9 plt.legend(fontsize = 18)
10 plt.tick_params(axis = 'x', labelsize = 18)
11 plt.tick_params(axis = 'y', labelsize = 18)
12 plt.show()

```

WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label




▼ Now subtract the moving average from that, and we apply ADF test.

```

1 data_log_minus_mean = data_adj_close_log - moving_average
2 data_log_minus_mean.dropna(inplace=True)
3 test_adf(data_log_minus_mean)

```




```
Results of Dicky Fuller Test
Test Statistics      -7.958585e+00
p-value             2.990744e-12
No. of lags used    2.300000e+01
Number of observations used  1.200000e+03
Critical value (1%)  -3.435811e+00
Critical value (5%)  -2.863952e+00
Critical value (10%) -2.568054e+00
dtype: float64
```

It is evident that test statistics is less than the critical values and $p < 0.05$, hence we reject the null hypothesis meaning our time series data is stationary.

▼ Train-Test Split

▼ We split our train test data by the ration of 75:25.

```
1 train_data, test_data = data_adj_close_log[:int(len(data_adj_close_log)*0.75)], data_adj_close_log[int(len(data_log_minus_mean)*0.75):
2 plt.figure(figsize = (20,10))
3 plt.xlabel('Dates', fontsize = 18)
4 plt.ylabel('Closing Prices', fontsize = 18)
5 plt.plot(data_adj_close_log, 'green', label = 'Train Data')
6 plt.plot(test_data, 'blue', label = 'Test Data')
7 plt.legend(fontsize = 18)
8 plt.tick_params(axis = 'x', labelsz = 18)
9 plt.tick_params(axis = 'y', labelsz = 18)
10 plt.savefig('train-test.png')
```



▼ Apply auto_arima function

- ▼ auto_arima function helps to find an optimal order for an ARIMA model. It returns the best ARIMA model

```

1 model_autoARIMA = auto_arima(train_data,
2                               start_p = 0, start_q = 0,
3                               test = 'adf',
4                               max_p = 7, max_q = 7,
5                               start_P = 0, D = 0,
6                               m = 1,
7                               d = None,
8                               seasonal = False,
9                               trace = True,
10                              error_action = 'ignore',
11                              suppress_warnings = True,
12                              stepwise = True
13                              )

```

```

Performing stepwise search to minimize aic
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-4762.672, Time=0.14 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=-4766.643, Time=0.33 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=-4767.001, Time=0.35 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-4762.192, Time=0.13 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=-4765.326, Time=0.39 sec
ARIMA(0,1,2)(0,0,0)[0] intercept : AIC=-4765.527, Time=0.56 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=-4763.580, Time=2.76 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=-4766.030, Time=0.11 sec

```

```

Best model: ARIMA(0,1,1)(0,0,0)[0] intercept
Total fit time: 4.804 seconds

```

```
1 model_autoARIMA.summary()
```

```

SARIMAX Results
Dep. Variable: y      No. Observations: 926
Model: SARIMAX(0, 1, 1) Log Likelihood 2386.501
Date: Thu, 03 Aug 2023 AIC -4767.001
Time: 04:17:26 BIC -4752.512
Sample: 0 HQIC -4761.473
- 926

Covariance Type: opg
coef std err z P>|z| [0.025 0.975]
intercept 0.0009 0.001 1.709 0.088 -0.000 0.002
ma.L1 -0.0824 0.023 -3.633 0.000 -0.127 -0.038
sigma2 0.0003 8.48e-06 39.641 0.000 0.000 0.000

Ljung-Box (L1) (Q): 0.00 Jarque-Bera (JB): 901.01
Prob(Q): 0.99 Prob(JB): 0.00
Heteroskedasticity (H): 2.58 Skew: 0.09
Prob(H) (two-sided): 0.00 Kurtosis: 7.83

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```

1 from statsmodels.tsa.arima.model import ARIMA
2
3 model = ARIMA(train_data, order=(1,1,0))
4 fitted = model.fit()

```

```

/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: A date index has been provided, but it
self._init_dates(dates, freq)
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: A date index has been provided, but it
self._init_dates(dates, freq)
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: A date index has been provided, but it
self._init_dates(dates, freq)

```

```

1 forecast = fitted.forecast(15, alpha=0.05)
2 print(forecast.values)

[6.88973349 6.88968309 6.88968699 6.88968669 6.88968671 6.88968671
 6.88968671 6.88968671 6.88968671 6.88968671 6.88968671 6.88968671
 6.88968671 6.88968671 6.88968671]
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:834: ValueWarning: No supported index is available. Predi
return get_prediction_index(

```

```

1 # prediction_series = pd.Series(prediction,index = test_data.index)
2 # fig, ax = plt.subplots(1, 1, figsize = (15, 5))
3 # ax.plot(data_adj_close_log)
4 # ax.plot(prediction_series)
5 # ax.fill_between(prediction_series.index,
6 #                 cf[0],
7 #                 cf[1],color='grey',alpha=.3)

```

```

1 prediction, confint = model_autoARIMA.predict(n_periods=test_data.shape[0], return_conf_int=True)
2
3 prediction
4 cf= pd.DataFrame(confint)

```

```

/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:834: ValueWarning: No supported index is available. Predi
return get_prediction_index(

```

```

1 fc, conf = model_autoARIMA.predict(n_periods = test_data.shape[0], return_conf_int = True)
2
3 fc_series = pd.Series(fc)
4
5 fc_series.index = test_data.index
6
7 lower_series = pd.Series(conf[:, 0], index = test_data.index)
8 upper_series = pd.Series(conf[:, 1], index = test_data.index)
9
10 plt.figure(figsize = (12, 5), dpi = 100)
11 plt.plot(train_data, label = 'training')
12 plt.plot(test_data, color = 'blue', label = 'Actual Stock Price')
13 plt.plot(fc_series, color = 'orange', label = 'Predicted Stock Price')
14 plt.fill_between(lower_series.index, lower_series, upper_series,
15                 color = 'k', alpha = 0.05)
16 plt.title('HCL Stock Price Prediction')
17 plt.xlabel('Time', fontsize = 12)
18 plt.ylabel('Actual Stock Price', fontsize = 12)
19 plt.legend(loc = 'upper left', fontsize = 12)
20 plt.tick_params(axis = 'x', labels = 12)
21 plt.tick_params(axis = 'y', labels = 12)
22 plt.savefig('Forecast.png')
23 plt.show()

```

```

/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:834: ValueWarning: No support
return get_prediction_index(

```



▼ Performance of the Model

- ▼ To evaluate the model performance, performance metrics like Mean Squared Error(MSE), Mean Absolute Error(MAE), Root Mean Squared Error(RMSE), Mean Absolute Percentage Error(MAPE) are used.

```
1 mse = mean_squared_error(test_data, fc)
2 mae = mean_absolute_error(test_data, fc)
3 rmse = math.sqrt(mse)
4 mape = np.mean(np.abs(fc - test_data)/ np.abs(test_data))
5
6 print('MSE: '+str(mse))
7 print('MAE: '+str(mae))
8 print('RMSE: '+str(rmse))
9 print('MAPE: '+str(mape))
```

```
MSE: 0.00803370918155261
MAE: 0.07811002552885458
RMSE: 0.08963096106565303
MAPE: 0.01120845696536436
```

From the above output it is clear that the model is performing really well on the test dataset.

✓ 0s completed at 12:17 AM

